

Data-Driven Fleet Operations: Demand Analytics and Smart Charging for Urban Ride-Hailing in Chicago

Advanced Analytics and Applications

Team 06



Authors:

Louis Schreier (7381065)

Daniel Schumacher (7409287)

Timo Ercole (7408459)

Contact: student@smail.uni-koeln.de

Supervisor: Univ.-Prof. Dr. Wolfgang Ketter

Co-Supervisor: Janik Muires

Chair of Information Systems for Sustainable Society
Faculty of Management, Economics and Social Sciences
University of Cologne

2026-07-26

Table of contents

1	Introduction	1
1.1	Business Context	1
1.2	Business Goal and Data Mining Goal	1
1.3	Datasets	1
2	Data Description	1
2.1	Taxi Trip Data	1
2.2	Weather Data	2
2.3	Point-of-Interest Data	3
3	Data Preparation	4
3.1	Taxi Data Cleaning	4
3.1.1	Missing values	4
3.1.2	Consistency Checks	4
3.1.3	Billing Discrepancy	5
3.1.4	Outlier Removal	5
3.1.5	Two Cleaning Levels	5
3.2	Weather Data Preparation	6
3.3	Point-of-Interest Data Preparation	6
3.4	Feature Engineering	6
3.4.1	Spatial and Temporal Granularity	6
3.4.2	Features Used in the Forecasting Models	6
4	Descriptive Analytics	7
4.1	Temporal and Trip Patterns	7
4.2	Spatial Analytics	7
4.2.1	Spatial Discretization: H3 Hexagonal Grid	7
4.2.2	Trip Metric Distributions	8
4.2.3	Spatial Demand Analysis	9
4.2.4	Demand by Community Area and Hexagon Resolution	9
4.2.5	Temporal Demand Patterns Across Space	9
4.2.6	Temporal Demand Patterns	10
4.2.7	Heatmap: Hour of Day $\times$ Day of Week	10
4.2.8	Origin–Destination Analysis	11
4.2.9	Spatio-Temporal Net Flow Analysis	12
5	Data Analytics	13
5.1	Predictive Analytics	13
5.1.1	Problem Formulation and Validation Strategy	13
5.1.2	Support Vector Machine	13
5.1.3	Feedforward Neural Network	16
5.1.4	Model Comparison	18
5.2	Smart Charging	18
6	Conclusions and Recommendations	18
6.1	What the Analysis Shows	18
6.2	Recommendations for the Fleet Operator	19
6.3	Private or Public Charging	19
6.4	Limitations	20
6.5	Outlook	20

List of Figures

1	H3 hexagonal grid overlaid on Chicago community areas at resolutions 6, 7, and 8.	8
2	Distribution of trip duration, distance, fare, tips, total cost, and inter-trip idle time across the full dataset.	8
3	Choropleth maps of pickup demand aggregated at community area level and H3 resolutions 6, 7, and 8.	9
4	Pickup demand by 4-hour time bucket at H3 resolution 6, covering the full 24-hour cycle.	10
5	Pickup demand heatmap cross-tabulated by hour of day and day of week.	11
6	Origin–destination matrix for the top 10 community areas by total trip volume. .	11
7	Net flow (pickups-dropoffs) by H3 hexagon, aggregated across the full observation period.	12
8	Training and validation loss curves for the final neural network model.	17

1 Introduction

1.1 Business Context

A German automotive manufacturer is launching a ride-hailing platform for US customers, operated entirely by an electrified vehicle fleet. It's entering a mature market with no operational track record, so it needs data-driven guidance on how to manage electric taxi fleets in cities.

Bane & Wayne Partners (BWP) was engaged to provide that guidance. This report presents the work of BWP's Data Science Team, using historical Chicago ride-hailing data to characterize demand patterns, forecast it, and design a charging strategy for electric vehicles.

1.2 Business Goal and Data Mining Goal

The goal is to give the client actionable intelligence for launching and running an electric fleet in a major US city. The analysis answers three questions:

1. **Where and when is demand concentrated?** This drives fleet positioning, driver scheduling, and market entry prioritization.
2. **How accurately can future demand be forecast?** Better predictions cut idle time, lower per-trip energy use, and raise driver earnings.
3. **How should charging be managed to minimize cost?** Charging timed to electricity prices and predicted demand can cut operating cost sharply.

The work is split into four tasks. Task 1 cleans the raw trips, discretizes them onto Uber H3 hexagons, and enriches them with weather data. Task 2 uses descriptive spatial analytics to surface demand hotspots and flow patterns. Task 3 trains two models (a Support Vector Machine and a feedforward neural network) to forecast demand per hexagon and time bucket. Task 4 develops a Reinforcement Learning agent for smart-charging policies under stochastic energy demand and dynamic pricing.

1.3 Datasets

Three sources feed the analysis: historical Chicago taxi trips (2024 onwards, from the Chicago Data Portal), hourly weather from Open-Meteo, and OpenStreetMap points of interest for urban context.

2 Data Description

Three data sources feed the analysis: Chicago taxi trip records, hourly weather, and OpenStreet Map point-of-interest. Together they cover demand, environmental context, and urban structure.

2.1 Taxi Trip Data

Downloaded from the Chicago Data Portal on 10 May 2026, the dataset holds every taxi trip from January 2024 to early May 2026: **14,985,679 trips** across 21 columns. Seven temporal features (date, hour, day of week, month, year, week, and whether a day is a weekend) are derived at load time, bringing the working total to 28.

Table 1 lists the key fields.

Table 1: Core fields of the Chicago taxi trip dataset.

Column	Type	Description
trip_id	string	SHA-1 hash; unique journey identifier
taxi_id	string	SHA-1 hash; anonymised vehicle identifier
trip_start_timestamp	datetime (UTC−6)	Rounded to the nearest 15 minutes
trip_end_timestamp	datetime (UTC−6)	Rounded to the nearest 15 minutes
trip_seconds	float	Recorded journey duration in seconds
trip_miles	float	Recorded journey distance in miles
pickup_census_tract	float	Census tract of pickup ($\approx 54\%$ missing)
dropoff_census_tract	float	Census tract of dropoff ($\approx 56\%$ missing)
pickup_community_area	float	Community area of pickup (0.05% missing)
dropoff_community_area	float	Community area of dropoff (7.2% missing)
fare	float	Metered fare in USD
tips	float	Tip amount in USD
trip_total	float	Total charge billed to the passenger
payment_type	string	Credit card, cash, mobile, etc.
company	string	Taxi company name
pickup_centroid_lat/lon	float	Centroid of the pickup census tract
dropoff_centroid_lat/lon	float	Centroid of the dropoff census tract

Three portal design choices matter downstream. Both IDs are SHA-1 hashes, so there’s no personally identifiable information. All timestamps are rounded to the nearest 15-minute interval for privacy. This caps temporal precision at 15-minute bins but doesn’t affect hourly or 4-hourly aggregations. And location is given only at the census tract level: with roughly 800 tracts in Chicago the spatial grain is coarser than GPS, but each tract’s centroid lat/lon supports spatial joins.

2.2 Weather Data

Hourly weather data was collected from Open-Meteo Historical Archive API (Open-Meteo, 2024) for the same date range as the taxi data (January 2024 - May 2026). Because Chicago spans roughly 600 km² and exhibits meaningful microclimatic variation between lakefront, the urban core, and the far northwest, a single weather station would be inadequate representative. We therefore determined three representative weather zones by applying K-means clustering ($k = 3$) to all taxi pickup coordinates, placing zone centroids where demand activity is the highest.

The resulting zones are centred on the downtown core (41.886°N, 87.629°W), the O’Hare area (41.979°N, 87.903°W), and the South Side (41.785°N, 87.750°W).

Hourly weather comes from the Open-Meteo Historical Archive API (Open-Meteo, 2024) over the same date range (January 2024 - May 2026). Chicago spans about 600 km² with real microclimatic variation between lakefront, the urban core, and the far northwest, so one weather station won’t represent it. We defined three representative weather zones by K-means clustering ($k = 3$) on all pickup coordinates, placing zone centroids where demand activity is densest: the downtown core (41.886°N, 87.629°W), O’Hare (41.979°N, 87.903°W), and the South Side (41.785°N, 87.750°W).

Table 2: Hourly weather variables fetched from Open-Meteo.

Variable	Unit	Description
temperature_2m	°C	Air temperature at 2 m height
apparent_temperature	°C	Feels-like temperature (wind chill / heat index)
precipitation	mm	Total precipitation
rain	mm	Liquid precipitation
snowfall	cm	Snowfall accumulation
snow_depth	cm	Snow depth on ground
windspeed_10m	km/h	Wind speed at 10 m
windgusts_10m	km/h	Wind gust speed at 10 m
cloud_cover	%	Fractional cloud cover

The dataset contains **61,335 observations** (~20,445 hours for each zone). Temperatures range from -26.8 °C to 36.3 °C, with appraent temperatures as low as -32.9 °C due to wind chill, which is consistent with Chicago’s continental climate. Median hourly precipitation is near zero while median cloud cover stands at 83 %, reflecting the city’s characteristically overcast conditions.

2.3 Point-of-Interest Data

POI data comes from OpenStreetMap via `osmnx` (Boeing, 2017), querying four tag categories for Chicago: `amenity`, `shop`, `tourism`, and `leisure`. The raw tags (e.g. restaurant, bus_stop, hospital) are too fine-grained to use directly, since many are too sparse for stable per-hexagon counts, and their demand signals overlaps (e.g. café and bar both mean nightlife). We collapsed them into seven demand-relevant categories, grouping tags with similar expected demand timing and dropping tags too niche to matter. After converting polygons to centroids and discarding unmapped types, **9,124 POIs** remain (Table 3).

Table 3: POI counts by functional category (OpenStreetMap, Chicago).

Category	Count	Examples
Food & Nightlife	5,640	Restaurants, bars, cafés, fast food
Education	1,219	Schools, universities, libraries
Shopping	1,084	Supermarkets, clothing, electronics
Healthcare	469	Hospitals, clinics, pharmacies
Entertainment	348	Theatres, cinemas, museums
Accommodation	229	Hotels, hostels

Category	Count	Examples
Transport	135	Bus stops, subway entrances, ferry terminals

POIs are static: they don't vary over time and aren't counted as demand. They are aggregated to H3 resolution 8 as per-category counts and joined into the feature matrix as a proxy for how much each cell attracts trips.

3 Data Preparation

3.1 Taxi Data Cleaning

Public portal data needs validation before use. Rather than applying fixed rules upfront, we worked through the taxi dataset step by step, by surfacing each quality issue, diagnosing its cause, and choosing a justified fix. Cleaning removes just over 10 % of raw records, retaining 89.99 % for analysis.

3.1.1 Missing values

Different fields are missing for different reasons, so each gets its own treatment (Table 4).

Table 4: Missing-value handling by field.

Field	Missing	Cause	Treatment
Timestamps	< 0.01 %	Fall in the autumn DST fold, where one hour repeats and the occurrence is ambiguous	Remove
Drop-off community area	7.2 %	Drop-off outside the 77 official community areas (suburban)	Recover via spatial join where inside the city; retain suburban trips unassigned
Trip duration	small	Not reported but derivable	Impute from end — start timestamp
Trip distance	small	Not reported	Estimate with great-circle distance (understates road distance)
Vehicle ID	8 records	Unattributable to any taxi	Remove

3.1.2 Consistency Checks

Beyond missing values, some records are present but aren't real journey. **Ghost trips.** have zero seconds, identical start/end time, identical start/end location, which look like a meter switched on and off. All four conditions must hold, so genuine short trips survive, so we remove the **153,078** that qualify (~1%). A further **1,014,515 trips** (6.8 %) log zero distance despite positive duration, with a median of 46 seconds and a quarter under 6 seconds. These meter

cancellations or GPS failures are all excluded. These are overwhelmingly meter cancellations or GPS failures rather than real journeys. Any trip with zero recorded distance is excluded. We also confirmed no end-before-start inversions and no duplicate trip IDs.

3.1.3 Billing Discrepancy

The billed total and the sum of its components (fare, tips, tolls, extras) disagree in 47.9 % of rows. The gap is always positive, almost always exactly \$0.50, and concentrated in credit-card and mobile payment transactions. This pattern is consistent with a card processing surcharge that's added to the total but is not itemised. We apply no correction and treat `trip_total` as the authoritative revenue figure.

3.1.4 Outlier Removal

After dropping non-journeys, we remove physically implausible values using two derived metrics (average speed and fare per mile). Table 5 lists the thresholds counts.

Table 5: Outlier removal criteria applied to the taxi trip dataset.

Criterion	Records removed	Share of dataset
Non-positive fare	5,243	0.04 %
Distance exceeding 100 miles	960	0.01 %
Duration exceeding 3 hours	11,629	0.08 %
Average speed above 70 mph	6,663	0.05 %
Average speed below 2 mph	155,478	1.13 %
Fare below \$2 per mile	28,537	0.21 %
Fare above \$100 per mile	small	< 0.01 %

As a check, correlation between per-neighbourhood pickup counts before and after outlier removal is 1.0000, meaning the spatial demand signal is fully preserved.

3.1.5 Two Cleaning Levels

Not every task requires the same level of cleaning. For **demand forecasting**, a trip with an odd speed or fare is still a real demand event. For **trip-level analyses** like duration or fare modelling, those records distort results. So we keep two versions:

Table 6: Rows retained under each cleaning level.

Level	What is removed	Rows retained
Demand (minimal)	DST-ambiguous timestamps	14,984,551 (99.99 %)
Trip (full)	Above, plus missing vehicle ID, ghost trips, zero-distance trips, and speed/fare outliers	13,485,728 (89.99 %)

Demand forecasting (Section 4-5) uses the minimal cleaning level; the profitability analysis uses the full level.

3.2 Weather Data Preparation

Zone assignment. The three weather zones from Section 2 come from K-means($k = 3$) on all pickup locations. Both the elbow method silhouette score pointed to $k = 3$.

**** Condition codes.**** The raw data carries a WMO condition code (0–99). Treating it as a number implies a false ordering (code 95 isn’t “95× worse” than code 1), and only three conditions actually appear: clear, rain, snow. Since rain and snow already have their own numeric columns, the code add nothing and is dropped.

Clock-change gaps. Each zone shows one hour gaps on spring-forward day (2024-03-10, 2025-03-09). The 02:00 hour doesn’t exist when clocks jump 01:00 → 03:00, so the gap is correct and no imputation is needed.

3.3 Point-of-Interest Data Preparation

The OSM query returns points, polygons, lines. All non-point features are reduced to their centroid, so every POI is a single coordinate. Each is mapped to one of seven functional categories via a lookup table, with unmapped types discarded. The final 9,124 POIs are assigned to the hexagon containing their coordinate.

3.4 Feature Engineering

The models don’t run on individual trips. Trips are aggregated to a grid: one row per (*hexagon, time bucket*), with pickup count as the target. Cells with no pickups are kept as explicit zero-demand rows, so the model learns from quiet periods too.

3.4.1 Spatial and Temporal Granularity

Grain sets a trade-off: finer units capture local variation but leave more cells at or near zero, which is harder to forecast. Table 7 shows the effect.

Table 7: Zero-demand share across spatial and temporal granularities.

Spatial unit	Time bucket	Grid cells	Zero-demand share
H3 resolution 5 (large hexagons)	1 hour	163,392	5.8 %
H3 resolution 6	1 hour	531,024	26.8 %
H3 resolution 6	3 hours	177,008	14.8 %
H3 resolution 7	3 hours	—	52.0 %

We chose **H3 resolution 6 with 3-hour buckets**: 177,008 cells (27 hexagons × 6,809 time buckets) at a manageable 14.8 % zero-demand share. Resolution 7 pushes zero-demand past 50% without a matching gain in operational precision.

3.4.2 Features Used in the Forecasting Models

Each feature is knowable before the bucket begins:

- **Calendar context.** - hour, day of week, month, weekend flag, US public holiday flag. These carry the daily and weekly rhythms from the descriptive analysis.

- **Cyclic encoding** - calendar variables are circular (23:00 is next to 00:00, December next to January), so each becomes a sine-cosine pair rather than a plain integer that implies a false break.
- **Weather conditions** - the nine variables for the nearest zone during the bucket, known in advance from forecasts.
- **Spatial identity** - a binary indicator per hexagon, letting the model learn a cell-specific baseline without assuming how neighbouring cells relate.

Together these give 50 input variables per grid cell.

4 Descriptive Analytics

4.1 Temporal and Trip Patterns

Before examining spatial structure, it is useful to characterize how demand is distributed across time and what individual trips look like. Taxi demand in Chicago is far from uniformly distributed across the day or week: it follows predictable rhythms tied to work schedules, social activity, and weather. Understanding these rhythms is a prerequisite for fleet positioning and driver scheduling.

On weekdays, demand follows a clear bimodal pattern: a morning peak around 08:00 driven by commuter traffic and a stronger, broader evening peak between 16:00 and 19:00. Overnight demand (00:00-06:00) is uniformly low across all days of the week, making this window a natural candidate for vehicle charging or maintenance. Weekends exhibit a fundamentally different profile: the morning commuter spike is absent, demand builds gradually through the afternoon and peaks in the early evening, reflecting leisure rather than work-driven travel. Midweek (Tuesday to Thursday) consistently records the highest absolute trip volumes.

4.2 Spatial Analytics

The spatial analysis characterizes where demand originates and how it is distributed across Chicago's geography. We work from coarse neighbourhood-level aggregations down to fine-grained hexagonal cells, and complement static demand maps with flow analysis and profitability metrics.

4.2.1 Spatial Discretization: H3 Hexagonal Grid

To enable a consistent and scalable spatial analysis, Chicago's geography was discretized using Uber's H3 hierarchical hexagonal grid system at resolutions 6, 7, and 8. Unlike census tracts, which follow irregular administrative boundaries, hexagons offer uniform area coverage and adjacency properties that are well-suited for spatial modeling.




Figure 1: H3 hexagonal grid overlaid on Chicago community areas at resolutions 6, 7, and 8.

Higher resolutions increase the number of hexagons exponentially while reducing their individual area, enabling finer-grained demand localization at the cost of data sparsity per cell. Resolution 7 was chosen as the primary unit for the subsequent analysis as it strikes a balance between spatial granularity and sufficient trip counts per hexagon for stable estimates.

4.2.2 Trip Metric Distributions

Before diving into spatial patterns, we first characterize the overall distribution of key trip metrics. All distributions are clipped at the 99th percentile to prevent outliers, such as airport flat-rate rides or exceptional idle gaps from distorting the visualization.




Figure 2: Distribution of trip duration, distance, fare, tips, total cost, and inter-trip idle time across the full dataset.

Several patterns stand out. Trip durations are right-skewed with a median around 600–700 seconds (~10 minutes), indicating that the majority of rides are relatively short urban hops. Trip

distances mirror this pattern, with most trips under 5 miles. Fare and total cost distributions confirm the dominance of low-cost short rides, with a long tail driven by airport and cross-city trips. The tip distribution has some extraordinary high counts at specific tip amounts, which could reflect the large share of cashless predefined tips. Idle time between consecutive trips per driver is concentrated in the 0–30 minute range, suggesting reasonably efficient utilization, though the mean around 60 minutes indicates periods of deliberate waiting or repositioning.

4.2.3 Spatial Demand Analysis

4.2.4 Demand by Community Area and Hexagon Resolution

To understand how spatial aggregation affects observed demand patterns, we compare pickup counts across three levels: community area, H3 resolution 6, and H3 resolutions 7 and 8.




Figure 3: Choropleth maps of pickup demand aggregated at community area level and H3 resolutions 6, 7, and 8.

At the community area level, the Near North Side and the Loop stand out as the dominant demand centers, consistent with their role as Chicago’s central business and entertainment district. The pattern is also visible in the airport community areas particularly O’Hare in the northwest, which generates a structurally distinct demand cluster far from the city center.

As resolution increases from H3-6 to H3-8, the hotspots sharpen considerably. At resolution 7, the spatial signal localizes to specific corridors: the Magnificent Mile, the River North entertainment district, and the area immediately surrounding O’Hare. Zones with moderate activity at coarser resolutions, such as the South Loop or Wicker Park, reveal internal heterogeneity at finer scales, with some sub-areas generating substantially higher demand than their neighbors. This demonstrates that aggregation at the community area level masks actionable sub-neighborhood variation that is relevant for fleet positioning decisions.

4.2.5 Temporal Demand Patterns Across Space

Demand does not only vary spatially but also shifts throughout the day. To capture this, pickups were aggregated into 4-hour time buckets at H3 resolution 7, and a choropleth map was produced for each bucket.




Figure 4: Pickup demand by 4-hour time bucket at H3 resolution 6, covering the full 24-hour cycle.

The northeastern lakefront and downtown corridor is the dominant demand hub across all time buckets, confirming its role as a structurally persistent hotspot regardless of the time of day. During the night window (00:00–03:59), overall demand is low but the downtown core retains relative activity, consistent with late-night hospitality and entertainment traffic. The early morning bucket (04:00–07:59) reveals a distinct secondary hotspot around Midway Airport in the southwest, attributable to early departures. The morning and midday buckets (08:00–15:59) see demand spread more broadly across the eastern half of the city, reflecting a mix of commuter, business, and leisure travel. The afternoon and evening windows (12:00–19:59) represent the peak demand period, with the downtown core intensifying sharply while outer neighborhoods remain comparatively quiet. From 20:00 onwards, activity tapers off but the lakefront strip retains a persistent hotspot driven by restaurant and nightlife traffic.

4.2.6 Temporal Demand Patterns

4.2.7 Heatmap: Hour of Day $\times$ Day of Week

The following heatmap cross-tabulates total pickups by hour of day and day of week, providing a compact view of how demand is structured across the weekly cycle.




Figure 5: Pickup demand heatmap cross-tabulated by hour of day and day of week.

The bulk of trips occur between 07:00 and 19:00 on weekdays, forming a dense demand block from Tuesday to Thursday that represents the highest observed demand in the dataset. Weekends show a distinctly different rhythm: demand starts later in the morning, lacks a pronounced commuter spike, and trails off more gradually in the evening, consistent with leisure rather than work-driven travel. The 00:00–06:00 window is uniformly quiet across all days, making it a natural candidate for driver rest periods or vehicle charging windows in an electric fleet context.

4.2.8 Origin–Destination Analysis

To understand mobility flows between areas, an origin–destination (OD) matrix was constructed for the top 10 community areas by outgoing trip volume.




Figure 6: Origin–destination matrix for the top 10 community areas by total trip volume.

The Near North Side and the Loop dominate both as origins and as destinations by a wide margin, reinforcing their central role in Chicago’s taxi network. Despite its geographic isolation in the northwest, O’Hare International Airport generates nearly 500,000 trips directed toward the Near North Side and Loop alone, confirming the airport-to-downtown corridor as one of the most important in the network. For a fleet operator, could imply that vehicles dropping off passengers in outer neighborhoods face a structural repositioning challenge to return to high-demand areas, in comparison to trips on the main demand corridor of the city.

4.2.9 Spatio-Temporal Net Flow Analysis

Beyond raw pickup and dropoff counts, net flow, defined as pickups minus dropoffs per hexagon, reveals where demand is structurally generated versus absorbed. This metric is directly actionable for fleet operators: zones with persistently positive net flow indicate areas where more vehicles are needed than are being returned by dropoffs, while negative net flow areas signal vehicle accumulation that may require active redistribution.




Figure 7: Net flow (pickups-dropoffs) by H3 hexagon, aggregated across the full observation period.

The downtown core and the O’Hare corridor show the most pronounced positive net flow, meaning that taxi demand in these areas consistently exceeds the volume of vehicles returned by dropoffs. The residential outskirts, by contrast, exhibit negative net flow, acting as vehicle sinks that absorb supply without generating equivalent demand in return. For an electric fleet operator, this spatial imbalance has direct implications for charging infrastructure placement: charging points situated in negative-flow residential areas could serve as productive holding

locations, where idle vehicles recharge between shifts as we expect idle times to be larger, and before being repositioned to high-demand corridors.

5 Data Analytics

5.1 Predictive Analytics

5.1.1 Problem Formulation and Validation Strategy

The forecasting task is to predict **taxi demand per spatial unit and time windows**, that is the number of pickups in a given hexagon during a given 3-hour period, using only information that is available before the window begins: the time of day, day of week, weather forecast, and location. The framing is cross-sectional rather than sequential. The goal is to answer questions such as “what demand can we expect on a rainy Sunday from 10:00-13:00 in the outskirts of the city?”. Not to project forward from the previous hour’s actual counts, which could not be available in a planning context.

Target transformation. Trip counts are heavily zero-inflated and right-skewed: a small number of downtown cells record hundreds of pickups per window while most outer cells record single digits. We apply a logarithmic transformation to the target before model training to stabilize the variance and improve model fit. All reported performance metrics are computed on the back-transformed, original count scale so that they are directly interpretable in terms of trips.

Train / validation / test split. The demand grid is randomly shuffled and divided 50 / 20 / 30 into training, validation, and test sets. The validation set is used exclusively for selecting model configurations, while the test set is reserved and touched only once, for the final reported metrics. Random shuffling is appropriate here because the models do not use lagged demand counts, there is no temporal information the model could inadvertently learn from the order of rows.

Baselines. Two naive benchmarks are evaluated alongside each model, providing context for the model results:

- *Per-cell mean:* for each hexagon and time window, the prediction is simply the historical average demand observed during training. $MAE = 26.77$ trips.
- *Climatology:* for each combination of hexagon, hour of day, and weekday-vs-weekend, the prediction is the historical average. This baseline captures the dominant daily and weekly seasonal rhythm. $MAE = 11.11$ trips.

Any model that fails to beat the per-cell mean offers no practical value. Beating climatology is the meaningful hurdle: that baseline already exploits the predictable weekly cycle that constitutes the bulk of foreseeable demand variation.

5.1.2 Support Vector Machine

Kernel Progression and Hyperparameter Tuning

A Support Vector Machine for regression (SVR) works by finding a function that stays within a tolerance band around the observed demand values while being as smooth as possible. The shape of that function is controlled by the kernel: a linear kernel allows only straight-line relationships, a polynomial kernel allows curved ones, and a Radial Basis Function (RBF) kernel allows highly flexible, non-linear relationships. Following the prescribed approach, we start from the simplest model and progressively increase complexity.

Three further settings are tuned simultaneously:

- **C** (regularisation strength): higher values allow the model to fit the training data more closely at the risk of generalising poorly.
- **Epsilon** (tolerance band width): predictions within this distance of the true value incur no penalty.
- **Gamma** (RBF and polynomial only): controls how strongly each training observation influences the model's prediction in its local neighbourhood.

Because the SVR becomes computationally demanding for large training sets, hyperparameter tuning was conducted on a random subsample of 10,000 training observations, with the best configuration then refit on up to 30,000 rows for final evaluation.

Kernel Comparison results

Table 8 shows the best result achieved per kernel family on the validation set.

Table 8: Best SVR configuration per kernel type, scored on the validation set.

Kernel type	Best regularisation	Hold-set MAE
Linear	$C = 1.0$	24.14 trips
Polynomial	$C = 1.0$	19.73 trips
RBF	$C = 10.0$	14.92 trips

The progression from linear (MAE 24.14) to RBF (MAE 14.92) illustrates the value of increasing model flexibility: the relationship between weather and calendar features and taxi demand is non-linear, and the RBF kernel's ability to model these curvatures brings a meaningful accuracy improvement. The best overall configuration (RBF kernel, $C = 10$, $\epsilon = 0.1$) is selected for final evaluation.

Final Test-Set evaluation

Table 9 compares the best SVR against the two baselines on the held-out test set of 53,103 grid cells.

Table 9: SVR test-set performance compared to baselines.

Model	MAE	RMSE	R^2
Per-cell mean	26.77	113.91	0.521
Climatology	11.11	55.11	0.888
SVR (RBF)	~15.8	—	~0.86

The SVR comfortably outperforms the per-cell mean but **does not beat the climatology baseline**. In other words, simply knowing the average demand for a given hexagon at a given hour on a weekday versus weekend outperforms the SVR on this test set. This outcome is informative rather than unexpected: without access to recent actual demand counts, the model is essentially learning the same seasonal pattern that the climatology already encodes, but less efficiently because it must discover that pattern from general-purpose features rather than from a precomputed schedule.

Where the Error Concentrates

The headline error figure conceals important structure in where the model struggles most.

Table 10: SVR test-set error decomposed by demand level.

Demand level (trips per window)	MAE	Observations
0–1	0.24	41,994
1–5	3.21	2,462
5–20	6.86	2,733
20–100	22.41	3,173
100+	179.32	2,741

The model is highly accurate on low-demand cells, which make up the vast majority of observations. However, it systematically under-predicts in high-demand cells, precisely the downtown and airport hexagons where accurate forecasting has the greatest operational value. This pattern reflects a known limitation of log-transformed regression: extreme values are compressed during training, causing the model to underestimate large demand spikes when predictions are returned to the original scale.

Effect of Spatial and Temporal Resolution

The assignment requires assessing how model performance varies across different granularity choices. We test the already-tuned model configuration across a range of hexagon sizes and time-window widths.

Table 11: SVR performance across spatial and temporal resolutions.

Hexagon size	Time window	MAE	R ²
Large (resolution 5, 8 cells)	1 hour	15.22	0.720
Large (resolution 5, 8 cells)	3 hours	36.98	0.841
Medium (resolution 6, 26 cells)	1 hour	6.50	0.657
Medium (resolution 6, 26 cells)	3 hours	16.25	0.719

As the spatial unit becomes finer, the absolute prediction error (MAE) decreases because individual hexagons contain fewer trips on average, but the model’s ability to explain variation (R²) also decreases as the demand signal becomes noisier in smaller cells. Hexagon size resolution 6 with 3-hour windows represents a practical balance. Using census tracts instead of hexagons as the spatial unit would introduce approximately 800 spatial units versus 27 at resolution 6, dramatically increasing sparsity and likely degrading performance further.

Model Shortfalls and Improvement Levers

The SVR’s inability to beat the climatology baseline reveals three structural limitations worth addressing in a follow-up project. First, the model has no access to **recent actual demand counts**. In time-series demand models, the previous period’s actual demand is consistently the strongest single predictor, the model effectively learns that “demand at 18:00 tomorrow will be similar to demand at 18:00 today.” Adding these lag features would likely push performance well above the seasonal baseline. Second, the SVR is **computationally constrained**: fitting on the full training set is infeasible, so tuning is conducted on a 10,000-row subsample. The model has not seen all the information available in the data. Third, the one-hot hexagon encoding treats each cell as **completely independent**. A spatially-aware representation, for example, encoding proximity to the city centre, distance to the nearest transport hub, or the demand level of neighbouring hexagons, would allow the model to generalise across cells rather than learning each one in isolation.

5.1.3 Feedforward Neural Network

Architecture and Hyperparameter Search

We train a fully connected feedforward neural network (FFNN) on the same demand grid and feature sets as the SVR, enabling a direct performance comparison. The network takes the 50-variable feature vector as input, passes it through stacked layers of neurons, and produces a single demand estimate as output. The depth, width, and training settings are selected through a systematic search.

The following configurations are explored:

Setting	Options tested
Hidden layer sizes	[128, 64], [256, 128]
Dropout (regularisation)	0 %, 10 %, 20 %
Learning rate	0.1, 0.01, 0.001
Optimisation algorithm	SGD, Adam

Each candidate is trained for up to 200 iterations, with training stopped early if the validation error stops improving for 20 consecutive iterations. The best configuration is then retrained from scratch on the combined training and validation data for the final test.

Activation Function Comparison

Using the best configuration found in the search above, we additionally compare three neuron activation function, the choices that determine how each layer transforms input, while keeping all other settings fixed.

Table 12: Validation error by activation function for the best network architecture.

Activation function	Validation MAE
ReLU	11.91 trips
Tanh	28.30 trips
Sigmoid	43.51 trips

ReLU substantially outperforms the alternatives. Tanh and sigmoid suffer from a training problem known as gradient saturation: neurons in these functions tend to produce near-constant outputs for many inputs, which slows or prevents effective learning in deeper networks. The final model uses ReLU throughout.

Final Test-Set Evaluation

The best configuration (two hidden layers of 128 and 64 neurons, ReLU activation, no dropout, learning rate 0.1, SGD optimiser) is retrained on the full training and validation data and evaluated once on the held-out test set.




Figure 8: Training and validation loss curves for the final neural network model.

Table 13: Neural network test-set performance compared to the climatology baseline.

Model	MAE	RMSE	R ²
Climatology	14.62	72.63	0.899
Neural network	12.31	65.25	0.919

Unlike the SVR, the neural network **outperforms the climatology baseline** on all metrics: MAE 12.31 versus 14.62, and R² of 0.919 versus 0.899. This confirms that the neural network captures demand patterns that go beyond the repeating weekly seasonal cycle, in particular, non-linear interactions between weather conditions, time of day, and location that a fixed seasonal average cannot represent. The model’s residual tendency to slightly under-predict (−7.11 trips on average) indicates that extreme demand events in the busiest downtown cells are still not fully captured. These spikes are driven by factors not present in the feature set (unplanned events, sports fixtures, concerts) and would require dedicated event-calendar data to forecast reliably.

Illustrative Demand Forecasting

One direct application of the fitted model is a context-specific demand estimate. Queried for expected demand in an outer-city hexagon on a rainy Sunday in July, from 10:00 to 18:00, the model returns:

Hour	Predicted trips
10:00	3.7
11:00	5.8
12:00	6.9
13:00	7.3
14:00	7.2
15:00	6.6
16:00	4.6
17:00	2.9
18:00	1.5
Total (8-hour window)	46.5

The gradual build-up and decline follows the leisure weekend pattern identified in the descriptive analysis. This kind of on-demand forecast (for any combination of location, time, weather, and day type) is the core operational output the model delivers to the fleet operator.

5.1.4 Model Comparison

A direct numerical comparison is complicated by a difference in time-window width between the two notebooks (3-hour windows for the SVR, 4-hour windows for the FFNN), which changes the absolute demand per cell and therefore the absolute MAE. The meaningful comparison is therefore each model’s performance relative to its own climatology baseline:

Table 14: Each model’s performance relative to the climatology baseline under its own granularity setting.

Model	Setting	Climatology MAE	Model MAE	Improvement
SVR (RBF)	H3 res. 6, 3 h windows	11.11	~15.8	Worse (−42 %)
Neural network	H3 res. 6, 4 h windows	14.62	12.31	Better (+16 %)

The neural network consistently outperforms the SVR in terms of improvement over the seasonal baseline. The gap has two root causes. First, the neural network trains on the full dataset using mini-batch gradient descent, it scales efficiently with data size. The SVR, by contrast, must be trained on a small subsample because of computational constraints, leaving most of the information in the training set unused. Second, the neural network learns interactions between features jointly across all layers, while the SVR with a fixed kernel applies the same transformation to every input, it has less flexibility to adapt to the interplay between, say, rainy weather and late-evening demand in airport hexagons.

Is the Neural Network Worth the Added Complexity?

For this demand forecasting task, the answer is yes. The SVR is easier to train and interpret, but it fails to improve on what a simple seasonal schedule already achieves. The neural network’s ability to beat the climatology baseline represents a meaningful operational gain: the model detects demand anomalies driven by weather and interactions between features, not merely reproducing the average weekly rhythm. For a fleet operator making positioning decisions under changing weather or on atypical weekdays, this distinction matters, the seasonal average provides no signal about whether tomorrow’s rainy Monday will see higher or lower demand than a dry one.

5.2 Smart Charging

6 Conclusions and Recommendations

6.1 What the Analysis Shows

Taken together, a consistent picture of how ride-hailing demand behaves in Chicago and what an incoming electric fleet should do about it is painted. Demand is concentrated in space and structured in time. A small number of hexagons in the downtown core, the Magnificent Mile, and the O’Hare corridor account for a disproportionate share of all pickups, while most of the city sits close to zero for much of the day. On top of that spatial structure sits a predictable weekly rhythm: a bimodal weekday pattern with a morning commuter peak and a stronger evening peak, a later and flatter weekend profile, and a quiet overnight window from midnight to six that is common to every day of the week.

These patterns are the foundation for the two forecasting models and the charging agent. The feedforward neural network beat the seasonal climatology baseline (MAE 12.31 against 14.62, R^2 0.919 against 0.899), while the Support Vector Machine did not. The reinforcement learning agent learned a charging policy that reliably met the driver’s departure energy needs at a cost well below a naive charge-at-maximum strategy. The sections below translate these results into recommendations for the client and set out where the analysis stops short.

6.2 Recommendations for the Fleet Operator

Position vehicles for the corridors, not the map as a whole. The strongest and most repeatable finding is that demand lives in a handful of places. For a fleet entering the market, this argues for a phased rollout that starts in the downtown-to-lakefront corridor and the airport connections rather than an even spread across all 77 community areas. The origin-destination analysis reinforces this: O’Hare alone sends close to half a million trips a year toward the Near North Side and the Loop, so the airport-to-downtown axis is a dependable source of paid mileage from day one.

Use the net-flow signal to plan repositioning. Pickups and dropoffs do not balance locally. The downtown core and the airport run a persistent positive net flow, meaning they hand out more trips than they take back, while residential neighbourhoods on the edges absorb vehicles without generating return demand. A fleet that ignores this will accumulate idle cars in low-demand areas during the day. Building repositioning logic around the net-flow map, and treating the outer sinks as places to hold and charge rather than places to wait for fares, addresses this directly.

Forecast for planning, and know what the forecast cannot see. The neural network is worth the added complexity because it responds to weather and to atypical days in a way that a fixed weekly schedule cannot. It is the right tool for shift planning and for answering concrete questions such as expected demand in a given area on a rainy Sunday afternoon. What it does not capture are one-off spikes from concerts, sports fixtures, and conventions, since nothing in the feature set flags those events. The operator should treat model output as a strong baseline to be overridden by hand on known event days, not as a complete demand picture.

6.3 Private or Public Charging

The charging results and the demand geography point to a mixed answer rather than a clean one. The cost function penalises high charging power convexly, so charging slowly over a long window is much cheaper than charging in short bursts at full rate. The agent exploited exactly this, spreading energy across the available slots and keeping a modest buffer above expected demand instead of racing to a full battery. That economics favours private or depot charging for the bulk of a vehicle’s energy needs: it lines up with the quiet overnight window the descriptive analysis identified, it lets each car draw at a low, cheap rate for hours, and it can sit in the negative-net-flow residential zones where cars would otherwise be idle anyway.

Public fast charging works against this cost structure, because its whole value is delivering a lot of energy quickly, which is the expensive corner of the curve. Our recommendation is therefore to anchor the fleet on private or depot charging placed near where vehicles naturally end their shifts, and to use public fast charging only as an opportunistic top-up during operations, for a car that would otherwise have to leave a high-demand corridor to recharge. Committing entirely to public infrastructure would expose the operator to both higher energy costs and competition for chargers during peak hours.

6.4 Limitations

Several caveats bound how far these findings should be pushed. Location in the source data is reported at census-tract centroids rather than true GPS coordinates, so every hexagon assignment, the hotspot detection, and the flow maps inherit a coarse spatial granularity; fine within-tract patterns are invisible by construction. The forecasting models see no recent demand counts and no event calendar, which is the main reason both under-predict the largest downtown spikes and why the SVR could not improve on a seasonal average. The SVR faced an additional practical ceiling: it had to be tuned on a small subsample because it does not scale to the full grid, so it was never shown all the available data. Finally, the charging agent is trained against a stylised environment with a fixed cost coefficient and a normally distributed departure demand. Real electricity prices vary hour to hour and real energy needs depend on the next day's trips, so the learned policy is a proof of concept rather than a deployable controller.

6.5 Outlook

The most valuable next step is to add lagged demand to the forecasting models. In demand forecasting the previous period's actual count is usually the single strongest predictor, and its absence here is the clearest gap between our setup and a production system. A spatially aware encoding of the hexagons, one that captures distance to the centre, proximity to transport hubs, and neighbouring demand, would also let the models generalise across cells instead of learning each one in isolation. On the data side, an event calendar covering concerts, conventions, and sporting fixtures would target precisely the spikes the current models miss, and real-time or day-ahead electricity prices would let the charging agent optimise against the tariffs a fleet actually faces rather than a flat coefficient. Further analyses worth pursuing on the existing data include a profitability study that pairs the demand maps with fare-per-mile and idle-time patterns to find where drivers earn most efficiently, and an explicit repositioning model that turns the net-flow findings into concrete vehicle movements. Each of these builds directly on the groundwork laid here and would move the analysis closer to an operational decision-support system for the client.

References

- Boeing, G. (2017). OSMnx: New methods for acquiring, constructing, analyzing, and visualizing complex street networks. *Computers, Environment and Urban Systems*, 65, 126–139. <https://doi.org/10.1016/j.compenvurbsys.2017.05.004>
- Open-Meteo. (2024). Historical weather API [Accessed: 2026-05-10].